

---

# CIS 678 Machine Learning

Convolutional Neural Networks (CNNs)

---

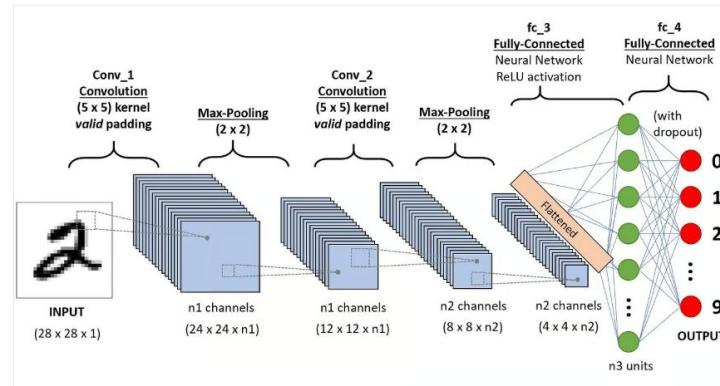
# Popular Neural Network (NN) Architectures

- *Feed Forward NNs*
- *Convolutional NNs*
- Sequence models
  - LSTM
  - Transformers
- Generative AI
  - GANs
  - VAEs
- Graph Neural Networks

---

# Popular Neural Network (NN) Architectures

- *Feed Forward NNs*
- **Convolutional NNs**
- Sequence models
  - LSTM
  - Transformers
- Generative AI
  - GANs
  - VAEs
- Graph Neural Networks





# Quiz!

---

# Question?

Q. Draw the diagram of a Feed Forward Neural Network with the properties given below, and estimate the minimum number of parameters your model would have:

1. **Input layer:** 2 nodes (to consume 2 input features,  $\{x_1, x_2\}$ )
2. **Two (2) Hidden layers** with the following configuration:
  - i) Hidden layer one: 2 nodes
  - ii) Hidden layer two: 2 nodes

**One bias input node** for each hidden layer in (2)
3. **Output layer:** 1 node (y)

---

# FFNN Parameters

$x_1$

$x_2$

Input layer

---

# FFNN Parameters

1

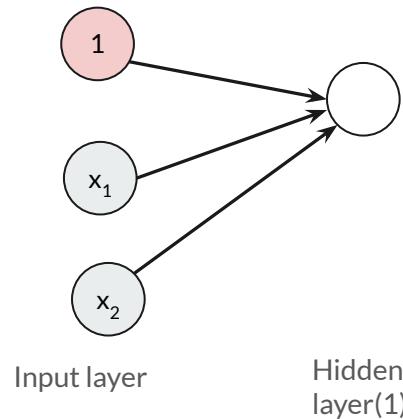
$x_1$

$x_2$

Input layer

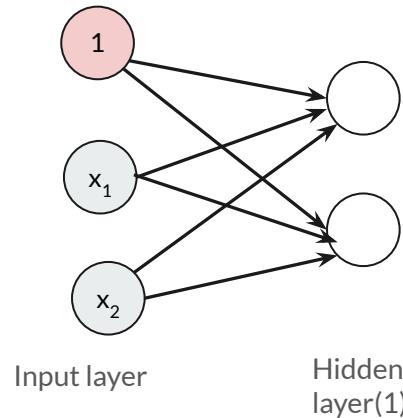
---

# FFNN Parameters



---

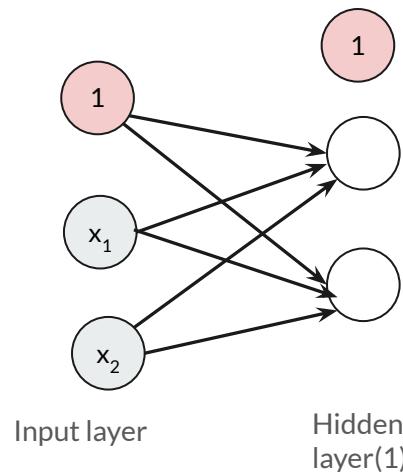
# FFNN Parameters



$$(2 + 1) \times 2$$

---

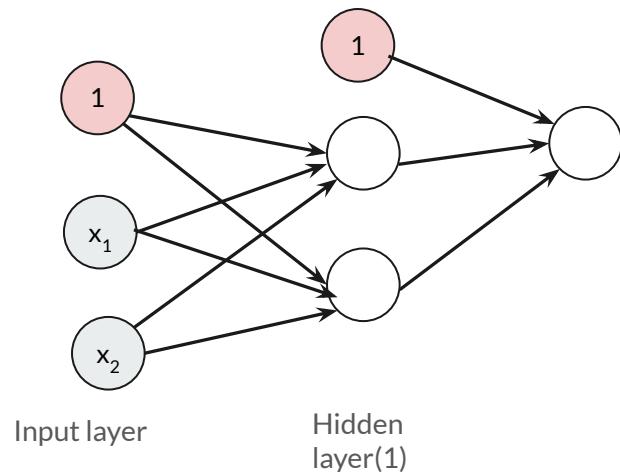
# FFNN Parameters



$$(2 + 1) \times 2$$

---

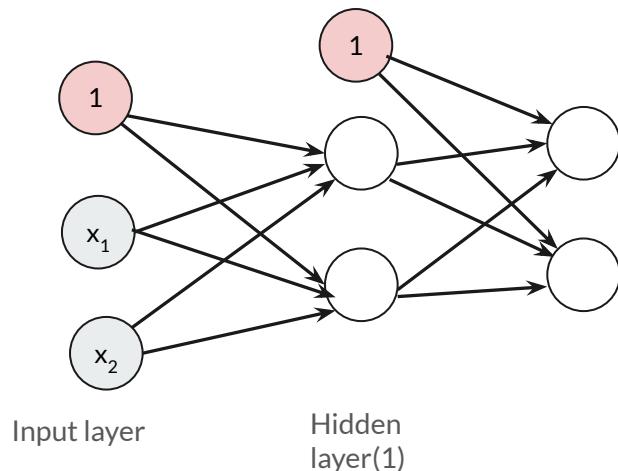
## FFNN Parameters



$$(2 + 1) \times 2 \quad +$$

---

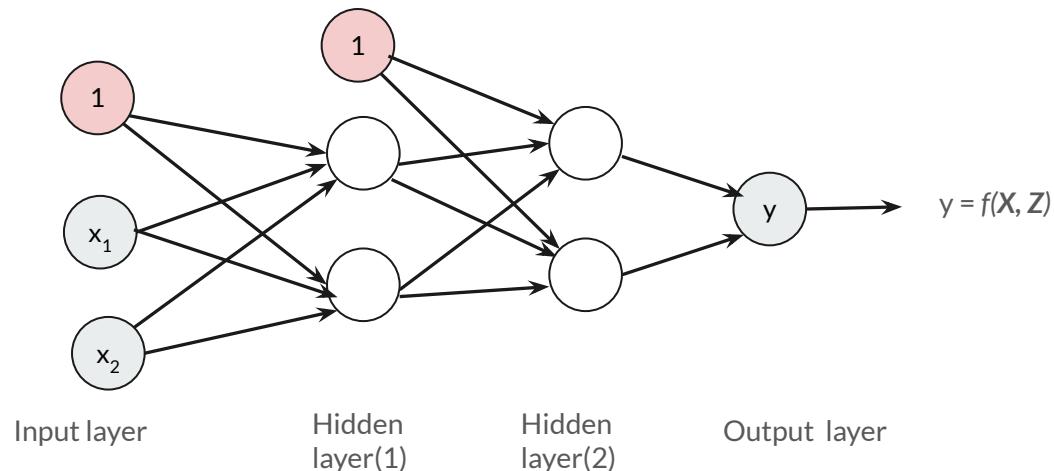
## FFNN Parameters



$$(2 + 1) \times 2 \quad + \quad (2 + 1) \times 2$$

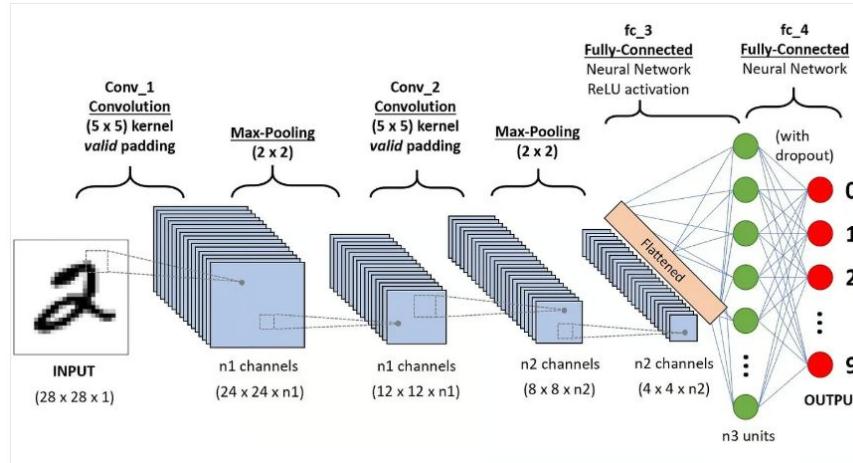
---

## FFNN Parameters



$$(2 + 1) \times 2 + (2 + 1) \times 2 + 2 \times 1 = 14$$

# Convolutional Neural Networks (CNNs)

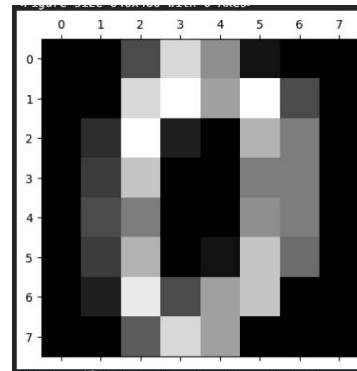


---

# Digit Classification examples

MNIST Image Classification

*What does it look like?*

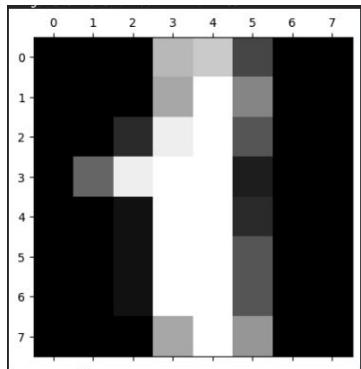
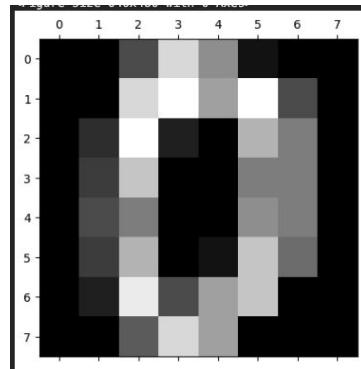




# Digit Classification examples

## MNIST Image Classification

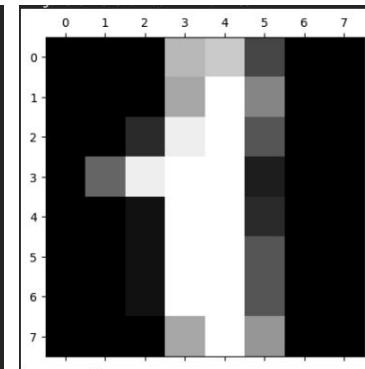
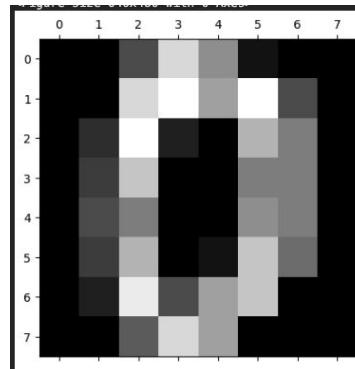
## *What does it look like?*



# Digit Classification examples

## MNIST Image Classification

Images of Size (2D Matrix): 8x8



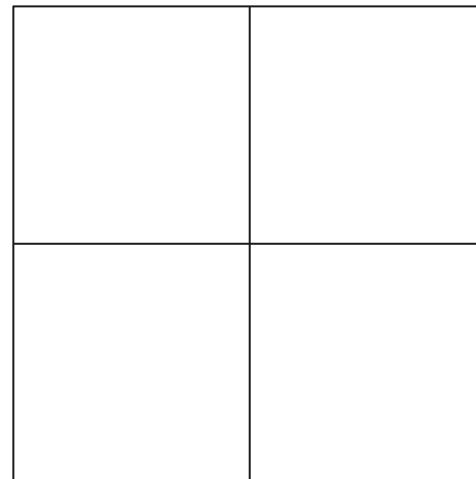
Corresponding Vector Representation of Size: 64

	pixel_0_0	pixel_0_1	pixel_0_2	pixel_0_3	pixel_0_4	pixel_0_5	pixel_0_6	pixel_0_7	pixel_1_0	pixel_1_1	...	pixel_6_6	pixel_6_7	pixel_7_0	pixel_7_1	pixel_7_2	pixel_7_3	pixel_7_4	pixel_7_5	pixel_7_6	pixel_7_7	
0	0.0	0.0	5.0	13.0	9.0	1.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	6.0	13.0	10.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	12.0	13.0	5.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	11.0	16.0	10.0	0.0	0.0	0.0
2	0.0	0.0	0.0	4.0	15.0	12.0	0.0	0.0	0.0	0.0	...	5.0	0.0	0.0	0.0	0.0	3.0	11.0	16.0	9.0	0.0	0.0
3	0.0	0.0	7.0	15.0	13.0	1.0	0.0	0.0	0.0	8.0	...	9.0	0.0	0.0	0.0	7.0	13.0	13.0	9.0	0.0	0.0	0.0
4	0.0	0.0	0.0	1.0	11.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	2.0	16.0	4.0	0.0	0.0	0.0

---

## Why FF Networks don't work well for Image Data?

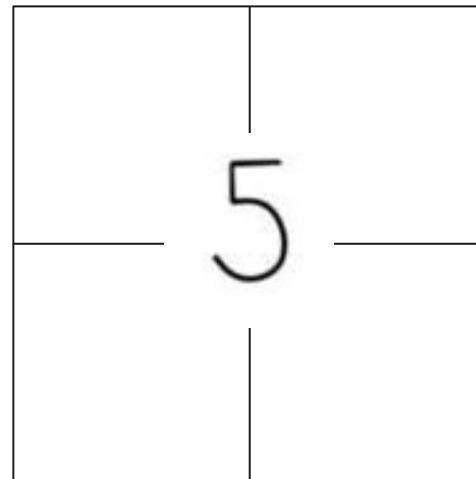
- We ask you to write a “5”



---

## Why FF Networks don't work well for Image Data?

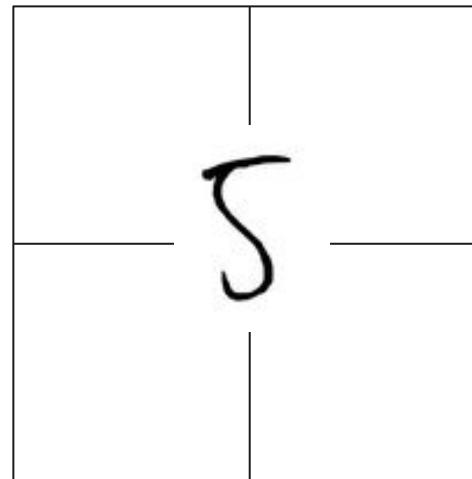
- We ask you to write a “5”
- A perfect “5” at the center of the coordinate



---

## Why FF Networks don't work well for Image Data?

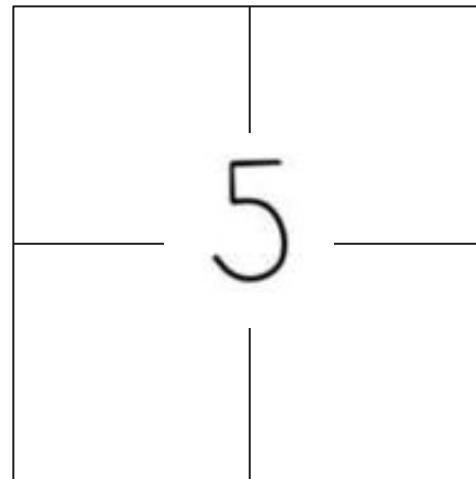
- We ask you to write a “5”
- A perfect “5” at the center of the coordinate
- This is my “5”, not good as yours but still at the center of the coordinate.



---

## Why FF Networks don't work well for Image Data?

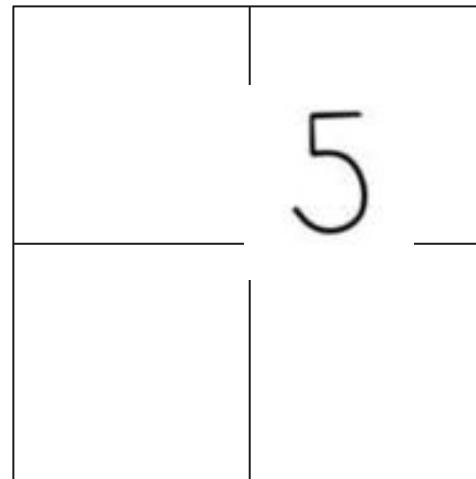
- We ask you to write a “5”
- A perfect “5” at the center of the coordinate
- Your “5” still will have a lots of variations



---

## Why FF Networks don't work well for Image Data?

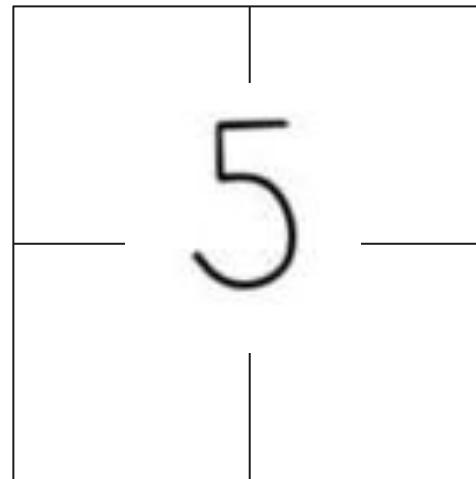
- We ask you to write a “5”
- A perfect “5” at the center of the coordinate
- Your “5” still will have a lots of variations
- Translated “5”



---

## Why FF Networks don't work well for Image Data?

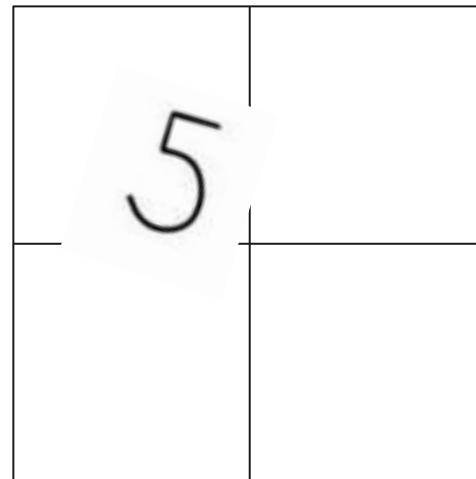
- We ask you to write a “5”
- A perfect “5” at the center of the coordinate
- Your “5” still will have a lots of variations
- Scaled “5”



---

## Why FF Networks don't work well for Image Data?

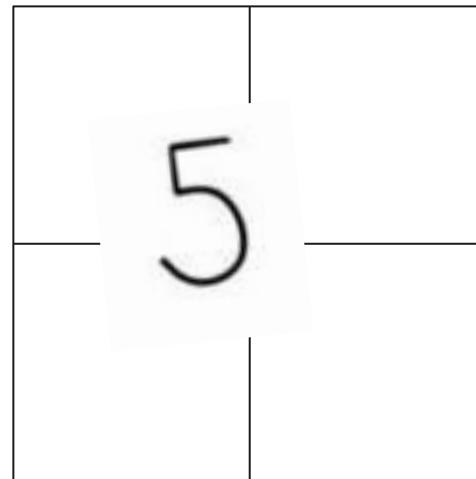
- We ask you to write a “5”
- A perfect “5” at the center of the coordinate
- Your “5” still will have a lots of variations
- Translated and Rotated “5”



---

## Why FF Networks don't work well for Image Data?

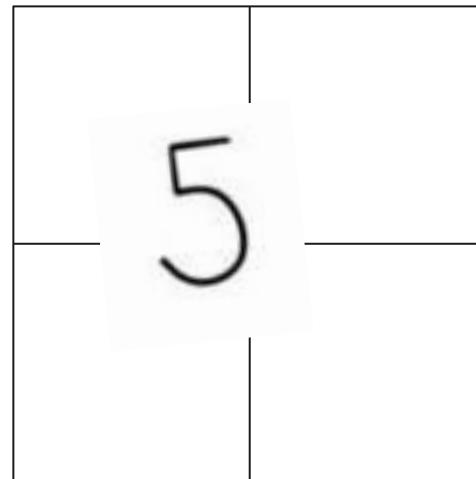
- We ask you to write a “5”
- A perfect “5” at the center of the coordinate
- Your “5” still will have a lots of variations
- Translated, Rotated and Scaled “5”



---

## Why FF Networks don't work well for Image Data?

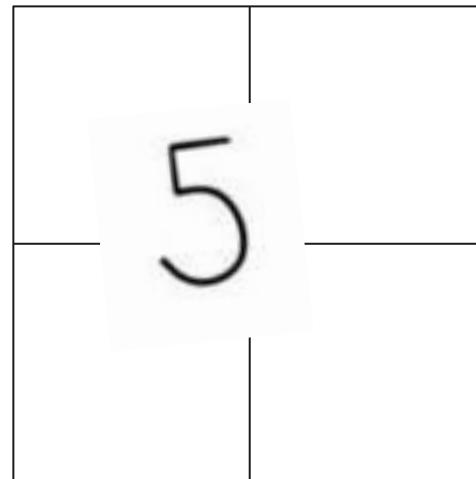
- We ask you to write a “5”
- A perfect “5” at the center of the coordinate
- Your “5” still will have a lots of variations and each definition is unique in the vector space



---

## Why FF Networks don't work well for Image Data?

- We ask you to write a “5”
- A perfect “5” at the center of the coordinate
- Your “5” still will have a lots of variations and each definition is unique in the vector space
- **How many different “5”s we will need provide the model to learn a perfect model?**



---

# Convolutional Neural Networks (CNNs)

- State of the Art for CV and some other problems
- Filters/Convolutional Kernels



# Convolutional NNs

- State of the Art for CV and some other problems
- Filters/Convolutional Kernels

Examples:

- *Alexnet*
- *VGGNet*
- *ResNet*
- *GoogLeNet*
- ..

---

# Convolution

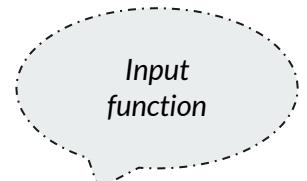
- Is a mathematical operation

$$(f * g)(t) := \int_{-\infty}^{\infty} f(\tau)g(t - \tau) d\tau.$$

---

# Convolution

- Is a mathematical operation

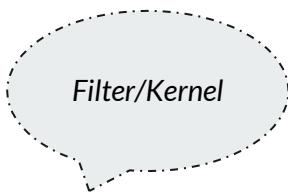

$$(f * g)(t) := \int_{-\infty}^{\infty} f(\tau) g(t - \tau) d\tau.$$

---

# Convolution

- Is a mathematical operation

$$(f * g)(t) := \int_{-\infty}^{\infty} f(\tau)g(t - \tau) d\tau.$$

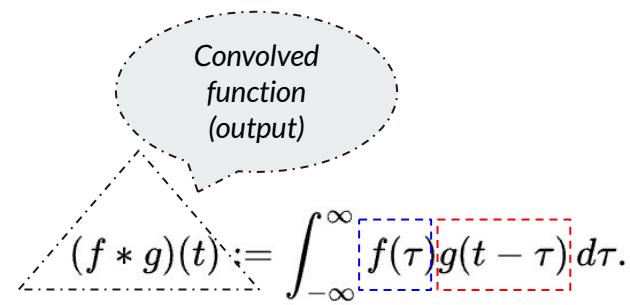


Filter/Kernel

---

# Convolution

- Is a mathematical operation



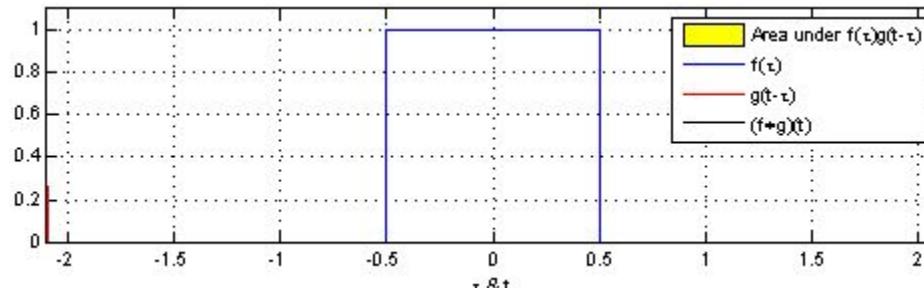
The diagram illustrates the convolution operation  $(f * g)(t)$ . A large dashed oval labeled "Convolved function (output)" contains a smaller dashed oval labeled " $f(\tau)g(t - \tau)$ ". Below the ovals is the convolution formula: 
$$(f * g)(t) := \int_{-\infty}^{\infty} f(\tau)g(t - \tau) d\tau.$$
 The term  $f(\tau)$  is highlighted with a blue dashed box, and the term  $g(t - \tau)$  is highlighted with a red dashed box, visually representing the weighted sum of the function values.

---

# Convolution

- Is a mathematical operation

$$(f * g)(t) := \int_{-\infty}^{\infty} f(\tau)g(t - \tau) d\tau.$$



Src

---

# Convolution

*Let's learn through a practical example*

## Moving average (Discrete case):

- Replaces each vector values pixel with an average of all the values in its neighborhood
- Moving average in 1D:

$$[1, 1, 1] / 3$$

$f(\tau)$

2	5	0	1	3
---	---	---	---	---

1D vector: 5 elements

ref

# Convolution

*Let's learn through a practical example*

## Moving average (Discrete case):

- Replaces each vector values pixel with an average of all the values in its neighborhood
- Moving average in 1D:

$$[1, 1, 1] / 3$$

$$g(t - \tau) d\tau.$$

1	1	1
---	---	---

1/3

*Known as Filter/Kernel*

$$f(\tau)$$

2	5	0	1	3
---	---	---	---	---

*1D vector: 5 elements*

ref

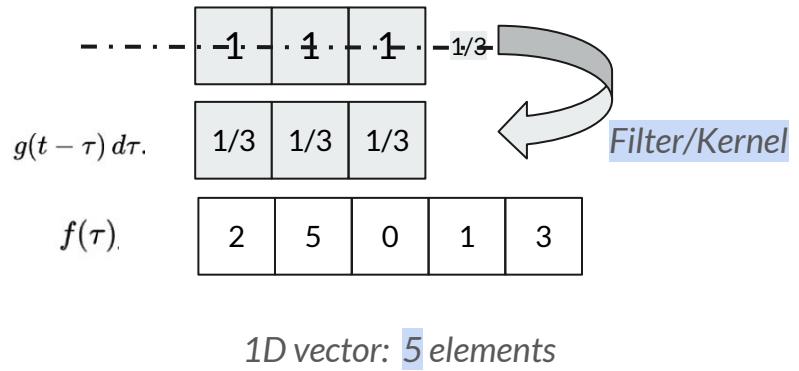
# Convolution

*Let's learn through a practical example*

## Moving average (Discrete case):

- Replaces each vector values pixel with an average of all the values in its neighborhood
- Moving average in 1D:

$$[1, 1, 1] / 3$$



ref

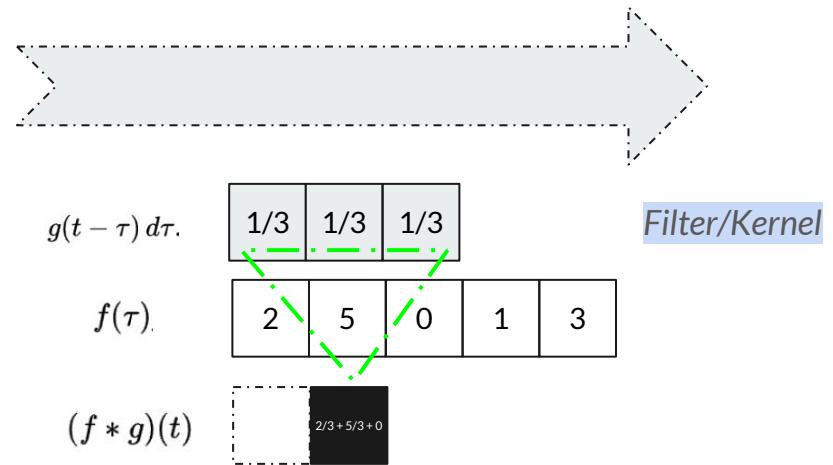
# Convolution

*Let's learn through a practical example*

## Moving average (Discrete case):

- Replaces each vector values pixel with an average of all the values in its neighborhood
- Moving average in 1D:

$$[1, 1, 1] / 3$$



ref

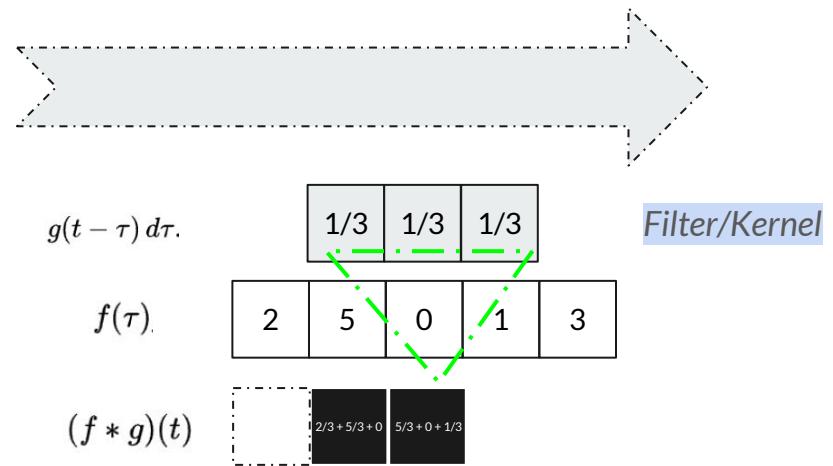
# Convolution

*Let's learn through a practical example*

## Moving average (Discrete case):

- Replaces each vector values pixel with an average of all the values in its neighborhood
- Moving average in 1D:

$$[1, 1, 1] / 3$$



ref

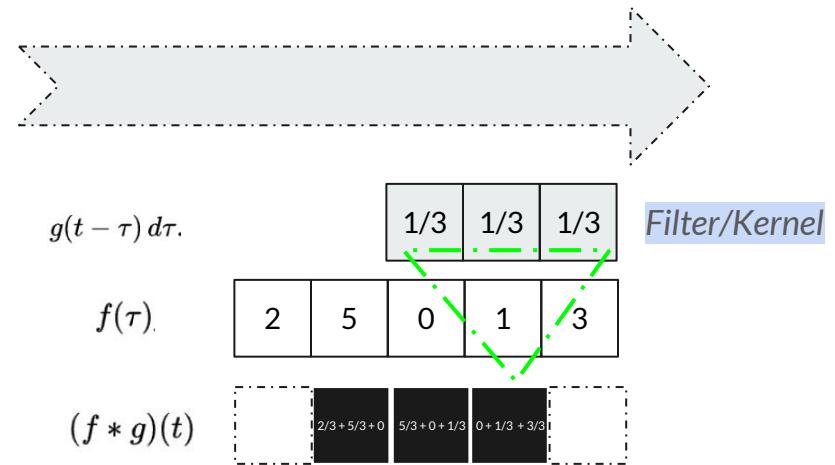
# Convolution

*Let's learn through a practical example*

## Moving average (Discrete case):

- Replaces each vector values pixel with an average of all the values in its neighborhood
- Moving average in 1D:

$$[1, 1, 1] / 3$$



ref

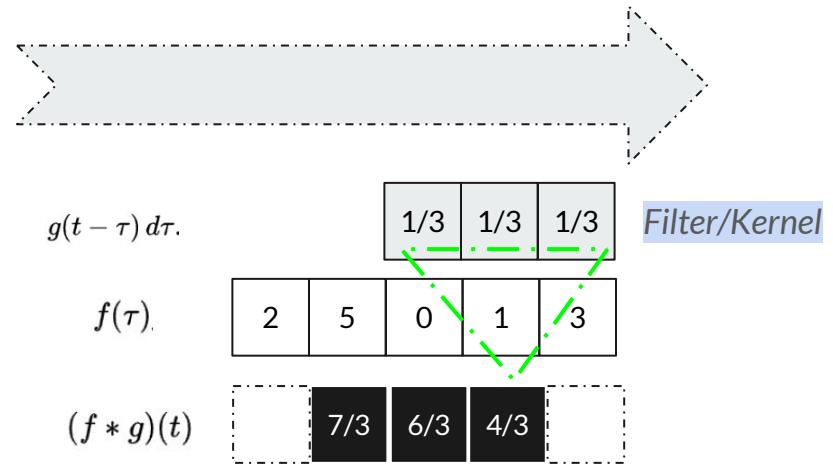
# Convolution

*Let's learn through a practical example*

## Moving average (Discrete case):

- Replaces each vector values pixel with an average of all the values in its neighborhood
- Moving average in 1D:

$$[1, 1, 1] / 3$$

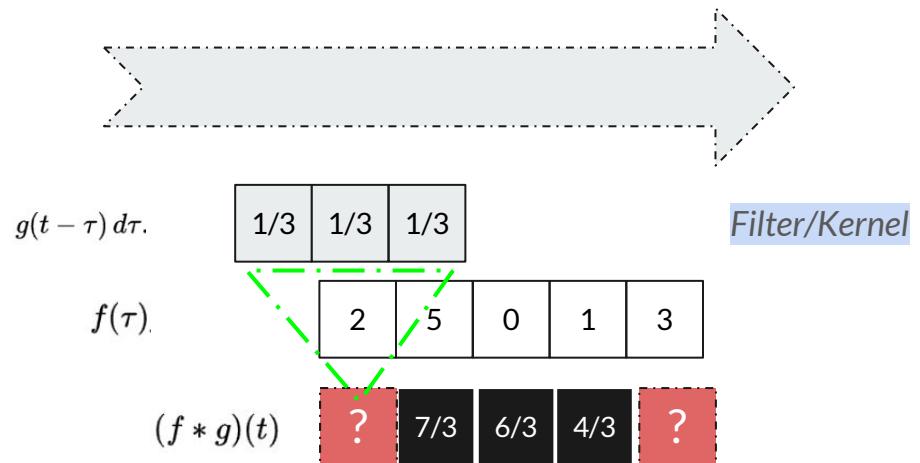


ref

# Convolution - Concepts

*How can we ensure that the input and output sizes are uniform?*

**Padding (0)**

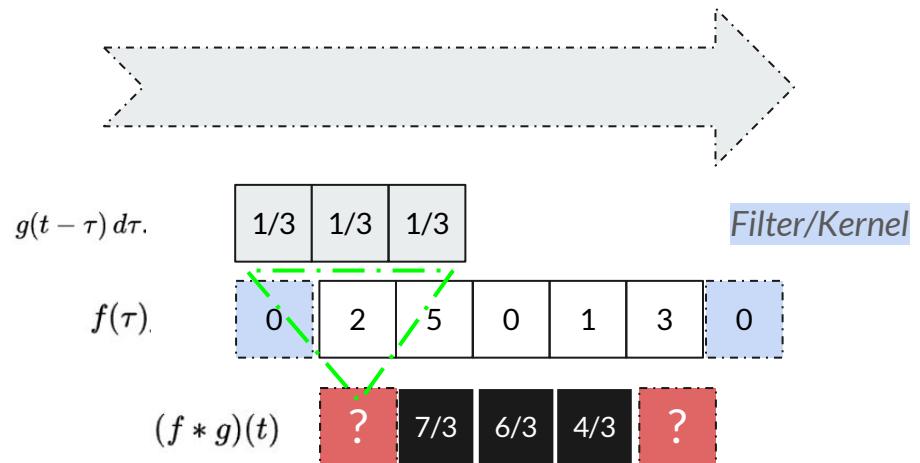


ref

# Convolution - Concepts

*How can we ensure that the input and output sizes are uniform?*

**Padding (0)**



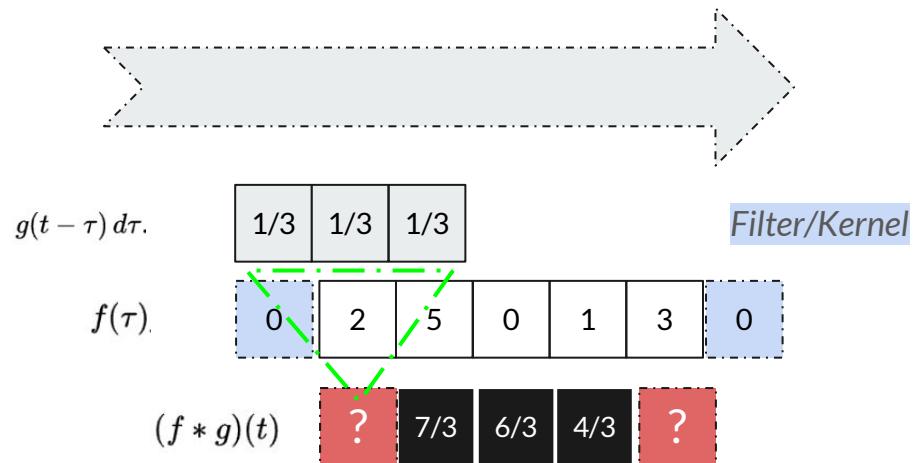
ref

# Convolution - Concepts

## **Faster Operation?**

## *Skip closely local context?*

## **Stride (1)**



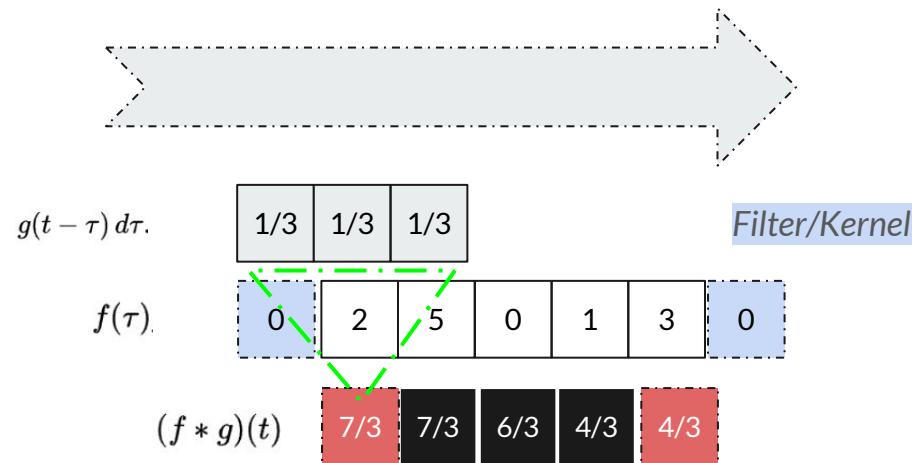
ref

# Convolution - Concepts

Faster Operation?

Skip closely local context?

Stride (1)



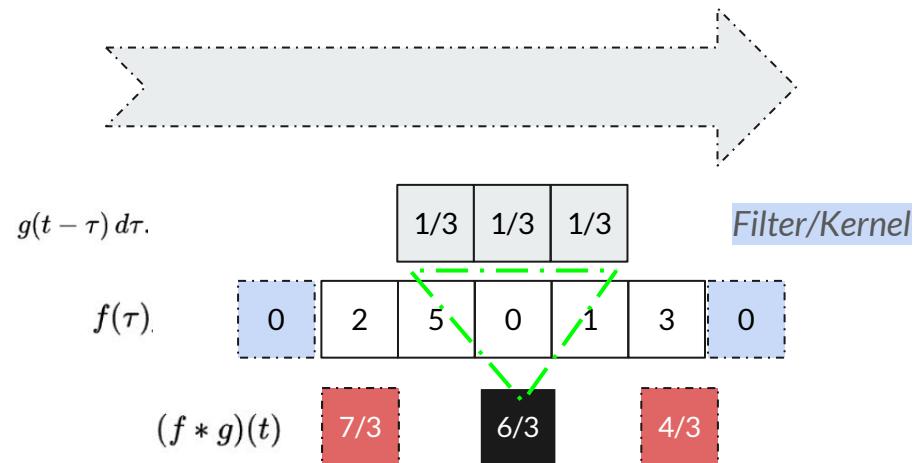
ref

# Convolution - Concepts

Faster Operation?

Skip closely local context?

Stride (2)



ref



## From 1D to 2D

*This is a general concept and applicable to  
any dimensions!*



# Image Filters

## Moving Average in 2D

- 2D convolution operation

Input image

$$F[x, y]$$

Filtered image

$$G[x, y]$$

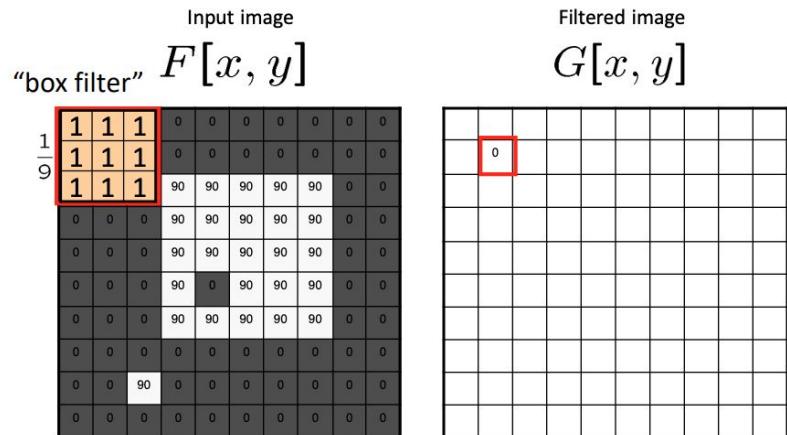
A large 10x10 grid of squares, used for drawing or plotting points.

---

# Image Filters

Moving Average in 2D

- 2D convolution operation



# Image Filters

## Moving Average in 2D

- 2D convolution operation

Input image

$$F[x, y]$$

Filtered image

$$G[x, y]$$

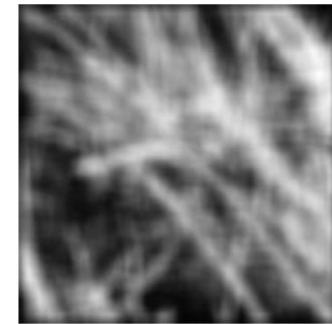
---

# Image Filters

## Moving Average in 2D

- 2D convolution operation
- Practical application (smoothing)

 ← Box filter:  
white = high value, black = low value



---

# Image Filters

## High Pass Filter

"If you take an image and blur it you only keep the "low frequencies". High pass makes the opposite, it only lets the "high frequencies" pass, or what most people call "**the details**". Any image can be deconstructed into these two components."[\[ref\]](#)

-1	-1	-1
-1	8	-1
-1	-1	-1



Note: Only for demonstration



# Image Filters

Other Filters:

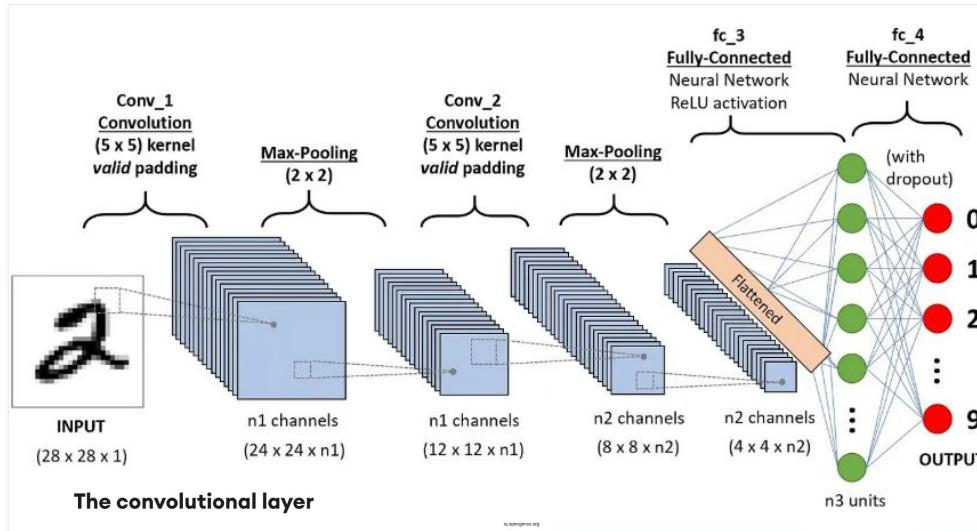
- Essentially these are convolution functions, **Prewitt filter** we also call as convolution kernels

$$G_x = \begin{bmatrix} -1 & 0 & +1 \\ -1 & 0 & +1 \\ -1 & 0 & +1 \end{bmatrix} \quad \text{and} \quad G_y = \begin{bmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ +1 & +1 & +1 \end{bmatrix}$$

**Sobel filter**

$$G_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix} \quad \text{and} \quad G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ +1 & +2 & +1 \end{bmatrix}$$

# Convolutional NNs



Input

1	0	1	0	1	0
0	1	1	0	1	1
1	0	1	0	1	0
1	0	1	1	1	0
0	1	1	0	1	1
1	0	1	0	1	0

Image patch (Local receptive field)

\*

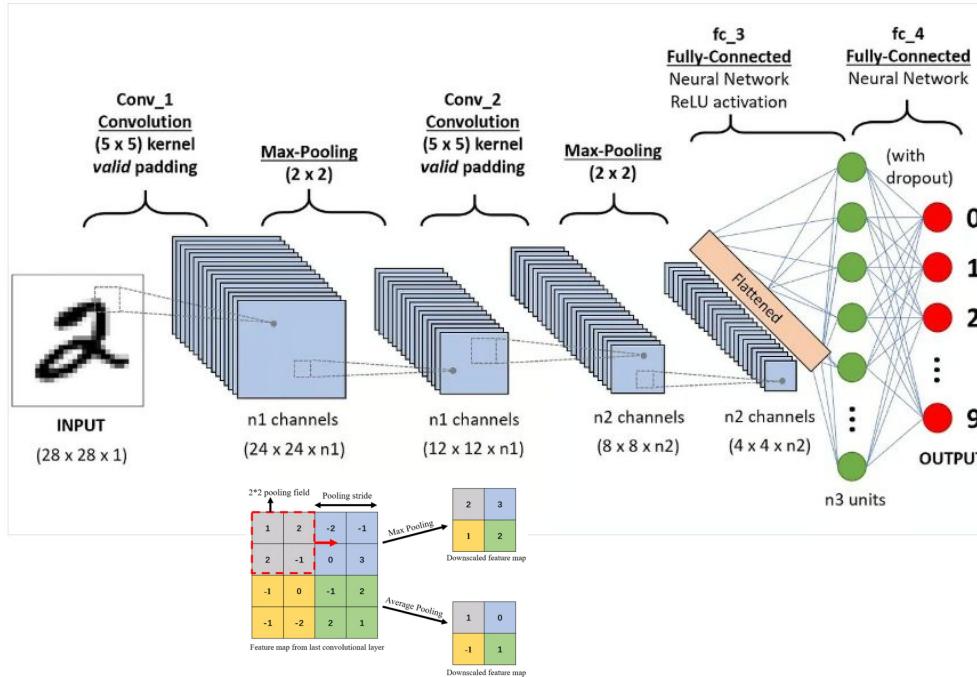
1	0	1	0	1	0
0	1	1	0	1	1
1	0	1	0	1	0
1	0	1	1	1	0
0	1	1	0	1	1
1	0	1	0	1	0

Kernel (filter)

Output

31
----

# Convolutional NNs





# Coding!

[Image classification \(Convolutional Neural Network\)](#)