
CIS 678 Machine Learning

Special Topics

Course Review Week(s)

Convolution

What will be the convolution output

No padding

$$g(t - \tau) d\tau.$$

0	1	0
---	---	---

Filter/Kernel

$$f(\tau)$$

2	5	0	1	3
---	---	---	---	---

$$(f * g)(t)$$

???

ref




Answer?

Convolution

What will be the convolution output

No padding




ref

Convolution

What will be the convolution output

No padding




ref

Convolution

What will be the convolution output

No padding



ref

2D Convolution

Data matrix (D):

1	0	1	1	0
1	0	0	1	1
0	1	0	0	0
1	0	0	0	1
0	1	0	1	0

Kernel matrix (D):

0	0	0
0	5	0
0	0	0



Answer?

2D Convolution

Data matrix (D):

1	0	0	1	0
0	0	0	1	1
0	0	0	0	0
1	0	0	0	1
0	1	0	1	0

Output Matrix

0		

2D Convolution

Data matrix (D):

1	0	0	0	0
1	0	5	0	1
0	1	0	0	0
1	0	0	0	1
0	1	0	1	0

Output Matrix

0	0	

2D Convolution


Data matrix (D):

1	0	1	0	0
1	0	0	1	0
0	1	0	0	0
1	0	0	0	1
0	1	0	1	0


Output Matrix

0	0	5


2D Convolution



Data matrix (D):



Output Matrix



The resulting 3x3 output matrix:

0	0	5
5		



Answer?

2D Convolution

Data matrix (D):

1	0	1	1	0
1	0	0	1	1
0	1	0	0	0
1	0	0	0	1
0	1	0	1	0

Kernel matrix (D):

0	0	0
0	5	0
0	0	0

Output Matrix

0	0	5
5	0	0
0	0	0



Our Plan

Our Plan


Vector space

Proximity or distance metric

- L1/Manhattan distance
- L2/Euclidean distance,
- Cosine distances

kNN model:

- Distance based
- Can be applied to both Regression and Classification tasks



Our Plan

Vector space

Proximity or distance metric


- L1/Manhattan distance
- L2/Euclidean distance,
- Cosine distances

kNN model:

- Distance based
- Can be applied to both Regression and Classification tasks

Probability (measuring uncertainty)

Probability distributions



Our Plan

Vector space

Proximity or distance metric

- L1/Manhattan distance
- L2/Euclidean distance,
- Cosine distances

kNN model:


- Distance based
- Can be applied to both Regression and Classification tasks

Probability (measuring uncertainty)


Probability distributions

ML introduction:

Linear to Polynomial Regression



Linear Regression



Model

$$\hat{y} = \beta_0 + \beta_1 x$$

$$\Theta = \{\beta_0, \beta_1\}$$

Fitting Error

$$\epsilon = |\hat{y} - y|$$

Optimization function

$$E_\Theta = \frac{1}{2} \sum_{i=1}^N (\hat{y}_i - y_i)^2$$

$$\Theta^* = \operatorname{argmin}_\Theta E\{(x_i, y_i)\}_{i=1, \dots, N}$$

Linear to Polynomial Regression




Table 1.1 Table of the coefficients w^* for polynomials of various order. Observe how the typical magnitude of the coefficients increases dramatically as the order of the polynomial increases.

	$M = 0$	$M = 1$	$M = 6$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*	-1.27	7.99	232.37	
w_2^*		-25.43	-5321.83	
w_3^*			17.37	48568.31
w_4^*				-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

Model generalization

Linear to Polynomial Regression




Table 1.1 Table of the coefficients w^* for polynomials of various order. Observe how the typical magnitude of the coefficients increases dramatically as the order of the polynomial increases.

	$M = 0$	$M = 1$	$M = 6$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*	-1.27	7.99	232.37	
w_2^*		-25.43	-5321.83	
w_3^*		17.37	48568.31	
w_4^*			-231639.30	
w_5^*			640042.26	
w_6^*			-1061800.52	
w_7^*			1042400.18	
w_8^*			-557682.99	
w_9^*			125201.43	

Absolute values
are increasing

Model generalization

Linear to Polynomial Regression

Model

$$\hat{y} = \beta_0 + \beta_1 x$$

$$\Theta = \{\beta_0, \beta_1\}$$

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$

Essentially, the same formulation

Generally ML vs Math conventions

$$W^* = \operatorname{argmin}_W E\{(x_i, t_i)\}_{i=1, \dots, N}$$

$$\epsilon = |\hat{y} - y|$$

Optimization function

$$E_\Theta = \frac{1}{2} \sum_{i=1}^N (\hat{y}_i - y_i)^2$$

$$\Theta^* = \operatorname{argmin}_\Theta E\{(x_i, y_i)\}_{i=1, \dots, N}$$

Linear to Polynomial Regression

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$

Regularizer

$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{1}{2} \|\mathbf{w}\|^2$$

Model generalization: Regularization

Table 1.1 Table of the coefficients w^* for polynomials of various order. Observe how the typical magnitude of the coefficients increases dramatically as the order of the polynomial increases.

	$M = 0$	$M = 1$	$M = 6$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*		-1.27	7.99	232.37
w_2^*			-25.43	-5321.83
w_3^*			17.37	48568.31
w_4^*				-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

Abs values
Are increasing

Linear to Polynomial Regression

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$

$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{1}{2} \|\mathbf{w}\|^2$$



Table 1.1 Table of the coefficients w^* for polynomials of various order. Observe how the typical magnitude of the coefficients increases dramatically as the order of the polynomial increases.

	$M = 0$	$M = 1$	$M = 6$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*		-1.27	7.99	232.37
w_2^*			-25.43	-5321.83
w_3^*			17.37	48568.31
w_4^*				-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

How to control this?

Model generalization: Regularization

decreasing

Linear to Polynomial Regression

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$




Table 1.1 Table of the coefficients w^* for polynomials of various order. Observe how the typical magnitude of the coefficients increases dramatically as the order of the polynomial increases.

	$M = 0$	$M = 1$	$M = 6$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*		-1.27	7.99	232.37
w_2^*			-25.43	-5321.83
w_3^*			17.37	48568.31
w_4^*				-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

Who to control this?



Probabilistic Twin

Of

The Least Squares Solution:


$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$

Probabilistic Twin

Of

The Least Squares Solution:

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$




Probabilistic Twin

Probabilistic Formulation: Modeling Error Distribution

$$p(\mathbf{t}|\mathbf{x}, \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(t_n | y(x_n, \mathbf{w}), \beta^{-1}).$$

Taking the log

$$\ln p(\mathbf{t}|\mathbf{x}, \mathbf{w}, \beta) = -\frac{\beta}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{N}{2} \ln \beta - \frac{N}{2} \ln(2\pi).$$



Probabilistic Twin

Probabilistic Formulation: Modeling Error Distribution

$$p(\mathbf{t}|\mathbf{x}, \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(t_n | y(x_n, \mathbf{w}), \beta^{-1}).$$


Taking the log

$$\ln p(\mathbf{t}|\mathbf{x}, \mathbf{w}, \beta) = -\frac{\beta}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{N}{2} \ln \beta - \frac{N}{2} \ln(2\pi).$$

Does it look familiar???


$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$

Maximizing Log likelihood is equivalent to minimizing the quadratic loss/error in the context of LR!





General ML models, Assignment 1,2, & Midterm






General ML models

Supervised

- | | |
|---|--|
| <ul style="list-style-type: none">• kNN• Linear Regression• Decision Tree• Meta learners<ul style="list-style-type: none">◦ Random Forest Regressor◦ Boosting Regressor• Support Vector Regressor (SVRs) | <ul style="list-style-type: none">• kNN• Logistic Regression• Decision Tree• Meta Learners<ul style="list-style-type: none">◦ Random Forest Classifier◦ Boosting Classifiers• Support Vector Classifiers (SVCs)• Naive Bayes |
| <h3>Regression</h3> | <h3>Classification</h3> |

General ML models





General ML models

Model generalization

- Universal concepts (applies to all models)
 - Cross validation
 - HP optimization
- Overfitting
- Underfitting

Universal concepts

Overfitting vs Under fitting

General ML models

Model generalization

- Training set, Validation set, Test set
- iid data

Training


Test /
Hold out set

- Overfitting
- Underfitting


Data splits

Overfitting vs Under fitting

General ML models



K-fold-cross validation



3-fold-cv


Train

validate

What HP gives the best validation score?

How to CV Sequential Data/Models

How to CV Sequential Data/Models




Sequential cross validation




Invalid configuration:
Training on the last two and validation on the **first fold**.

Sequential cross validation



Always
follow!

Sequential cross validation




*Always
follow!*

Sequential cross validation



Sequential cross validation



Always
follow!

Unsupervised Learning

Supervised Learning

- We learned about **Classification** and **Regression**
- These are examples of **supervised learning**
- In your data you have both **X(features)** and **y(Labels)**

$$f(y|X)$$

Classification

X	$y \in \{cat, dog, rabbit\}$
$x_1, x_2, \dots x_m$	cat
$x_1, x_2, \dots x_m$	rabbit
...	...
$x_1, x_2, \dots x_m$	dog

Label (y) is predefined

Regression

X	$y \in R$
$x_1, x_2, \dots x_n$	1.9
...	...
$x_1, x_2, \dots x_n$	2.5

Unsupervised Learning

- In contrast, in **Unsupervised learning**, we have to learn meaningful representations/models from **X(features)** only.
- **Clustering** is an example of unsupervised learning

X
$x_1, x_2, \dots x_o$
\dots
$x_1, x_2, \dots x_o$

No
concept
of data
label (y)

Unsupervised Learning

- In contrast, in **Unsupervised learning**, we have to learn meaningful representations/models from **X(features)** only.
- **Clustering** is an example of unsupervised learning


X
$x_1, x_2, \dots x_o$
...
$x_1, x_2, \dots x_o$

No
concept
of data
label (y)

$f(y|X)$


Clustering

- A medical study involving
 - Average number of symptoms per week
 - Average severity



Clustering

- A medical study involving
 - Average number of symptoms per week
 - Average severity



Usages of Clustering

Clustering has a myriad of uses in a variety of industries. Some common applications for clustering include the following:


- **Market segmentation**
- Social network analysis
- Search result grouping
- Medical imaging
- Image segmentation
- Anomaly detection



Usages of Clustering

Clustering has a myriad of uses in a variety of industries. Some common applications for clustering include the following:

- Market segmentation
- **Social network analysis**
- Search result grouping
- Medical imaging
- Image segmentation
- Anomaly detection




Clustering

- What is clustering?
- Clustering algorithms:
 - **K-Means:** Centroid Based
 - **Hierarchical clustering:** Distance connectivity based
 - **GMM:** Distribution based
 - **DBSCAN:** Density Based
- Identifying the number of clusters ?

Selecting the Number of clusters

- Elbow method
- Silhouette score based




Explained Variance. The "elbow" is indicated by the red circle. The number of clusters chosen should therefore be 4.

Curse of Dimensionality

Curse of Dimensionality

- We have added Weight as our second feature dimension
- Do you find the samples represent our class?
- What's the domain coverage now?
- Is ~10% a reasonable assumption?
- See how numbers dropped from ~40% to ~10%
- The other way, the empty space grew up from ~60% to ~90%
- **What's the number would look like if we add 2 more dimensions?**

- Empty space grows exponentially with the increase in adding new features.
- **Data distribution becomes sparse, and difficult to learn a good model.**



Dimensionality Reduction

General idea

- You have some data of feature dimension size, $|D|$
- You want to compress them to of size, $|d|$
- $|d| < |D|$




Dimensionality Reduction

- Linear
 - o Principal Component Analysis (PCA)
 - o SVD
- Neural Networks
 - o Auto Encoders
 - o RBMs


PCA

Notebook presentation


1 $C = \begin{bmatrix} 1 & 0 \\ 0 & 7 \end{bmatrix}$ $\lambda_{1,2} = [1 \ 7]$
 $V = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$




2 $C = \begin{bmatrix} 7 & 0 \\ 0 & 1 \end{bmatrix}$ $\lambda_{1,2} = [7 \ 1]$
 $V = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$



3 $C = \begin{bmatrix} 4 & 3 \\ 3 & 7 \end{bmatrix}$ $\lambda_{1,2} = [2.1 \ 8.8]$
 $V = \begin{bmatrix} -0.8 & -0.5 \\ 0.5 & -0.8 \end{bmatrix}$



4 $C = \begin{bmatrix} 4 & -3 \\ -3 & 7 \end{bmatrix}$ $\lambda_{1,2} = [2.1 \ 8.8]$
 $V = \begin{bmatrix} -0.8 & 0.5 \\ -0.5 & -0.8 \end{bmatrix}$




λ = eigenvalues

V = eigenvectors

What are we optimizing?


Two techniques (in general):

- Maximize variance
- Minimize errors (distance between each green-red pairs)



Unsupervised learning (nonlinear)

- Auto Encoders
- Restricted Boltzmann Machines (RBMs)




X

Z

Unsupervised learning (nonlinear)

- Auto Encoders
- Restricted Boltzmann Machines (RBMs)



X

Z

X'






Deep Learning

Assignment 3


Logistic Regression

- Probabilistic classifier
- Sigmoid function


$$p(x) = \frac{1}{1+e^{-(w_0+w_1x)}}$$




Neural Networks (Node)



Neural Networks (Node)




Neural Networks (Node)



Name	Plot	Function, $g(x)$
Identity		x
Binary step		$\begin{cases} 0 & \text{if } x < 0 \\ 1 & \text{if } x \geq 0 \end{cases}$
Logistic, sigmoid, or soft step		$\sigma(x) \doteq \frac{1}{1 + e^{-x}}$
Hyperbolic tangent (tanh)		$\tanh(x) \doteq \frac{e^x - e^{-x}}{e^x + e^{-x}}$
Soboleva modified hyperbolic tangent (smht)		$\text{smht}(x) \doteq \frac{e^{ax} - e^{-bx}}{e^{cx} + e^{-dx}}$
Rectified linear unit (ReLU) ^[13]		$(x)^+ \doteq \begin{cases} 0 & \text{if } x \leq 0 \\ x & \text{if } x > 0 \end{cases} = \max(0, x) = x\mathbf{1}_{x>0}$

A NN with ???
Activation function

Feed-forward (FF) neural networks



Error Back propagation


Gradient Descent (Error Back Propagation)

The Backwards Pass

Let's focus on $\frac{\partial E_{total}}{\partial w_5}$


What would be the gradient update for w5?

Adapted from



Gradient Descent (Error Back Propagation)

The Backwards Pass



$$E_{total} = \frac{1}{2}(target_{o1} - out_{o1})^2 + \frac{1}{2}(target_{o2} - out_{o2})^2$$


$$\frac{\partial E_{total}}{\partial out_{o1}} = 2 * \frac{1}{2}(target_{o1} - out_{o1})^{2-1} * -1 + 0$$

$$\frac{\partial E_{total}}{\partial out_{o1}} = -(target_{o1} - out_{o1}) = -(0.01 - 0.75136507) = 0.74136507$$

Adapted from

Gradient Descent (Error Back Propagation)

The Backwards Pass




$$out_{o1} = \frac{1}{1+e^{-net_{o1}}}$$

$$\begin{aligned}\frac{\partial out_{o1}}{\partial net_{o1}} &= out_{o1}(1 - out_{o1}) \\ &= 0.75136507(1 - 0.75136507) \\ &= 0.186815602\end{aligned}$$

Adapted from

Gradient Descent (Error Back Propagation)

The Backwards Pass





$$net_{o1} = w_5 * out_{h1} + w_6 * out_{h2} + b_2 * 1$$

$$\frac{\partial net_{o1}}{\partial w_5} = 1 * out_{h1} * w_5^{(1-1)} + 0 + 0 = out_{h1} = 0.593269992$$

Adapted from

Gradient Descent (Error Back Propagation)



Hidden Layer



- While changing w_5 affects only O_1 , a change in w_1 will change both O_1 and O_2


Gradient Descent (Error Back Propagation)

Hidden Layer



- Can you think of an arbitrary node in a giant and complex NN? What challenges we may encounter?

Vanishing and Exploding Gradients



$$o_1 = a_1(w_1 \cdot x)$$

$$o_2 = a_2(w_2 \cdot a_1(w_1 \cdot x))$$

..... (l=20)

[ref](#)




Convolutional NNs

- State of the Art for CV and some other problems
- Filters/Convolutional Kernels

Examples:

- *Alexnet*
- *VGG*
- *ResNet*
- *GoogLeNet*
- ..

Convolutional NNs



Input

1	0	1	0	1	0
0	1	1	0	1	1
1	0	1	0	1	0
1	0	1	1	1	0
0	1	1	0	1	1
1	0	1	0	1	0

Image patch (Local receptive field)

*


1	0	1	0	1	0
0	1	1	0	1	1
1	0	1	0	1	0
1	0	1	1	1	0
0	1	1	0	1	1
1	0	1	0	1	0

Kernel (filter)

Output

31

Convolutional NNs





How are the time-series problems different?

- The models (Regression and Classification), we have learned so far are of the form:

$$f(y|X)$$

- Purely time series models are of form :

$$f(y_t|y_{t-1}, y_{t-2}, \dots, y_0), \text{ where there is no explicit, } X.$$

- For certain data, especially the time series, we can take advantage of the form:

$$f(y_t|X, y_{t-1}, y_{t-2}, \dots, y_0); \text{ essentially, here the input features are } X \text{ plus the lagged instances of the target } y.$$


Statistical time series models

- Autoregressive Models (AR)
- Moving Average (MA)
- Autoregressive Moving Average (ARMA)
- Autoregressive Moving Integrated Average (ARIMA)

$$f(y_t | y_{t-1}, y_{t-2}, \dots, y_0)$$

Statistical time series models


- Autoregressive Models (AR)



$$f(y_t | y_{t-1}, y_{t-2}, \dots, y_0)$$

Statistical time series models


- Autoregressive Models (AR)



$$f(y_t | y_{t-1}, y_{t-2}, \dots, y_0)$$


Statistical time series models

- Autoregressive Models (AR)




$$f(y_t | y_{t-1}, y_{t-2}, \dots, y_0)$$

How to CV Sequential Data/Models






Final Project + Special Topics



Recurrent Neural Networks

Originally designed to Solve:

- Sequential Problems
 - o NLP (MT, ...), Speech Recognition,...
 - o DNA Sequencing
- Time-series Problems




ref

Recurrent Neural Networks

Originally designed to Solve:

- Sequential Problems
 - o NLP (MT, ...), Speech Recognition,...
 - o DNA Sequencing
- Time-series Problems




[ref](#)

Transformers

Examples:

- *Encoder decoder pair*
- GPT
- BERT

BERT
Encoder



[ref](#)




Generative AI

Examples:

- Generative Adversarial Networks (GANs)
- Variational Autoencoders (VAEs)

Generative AI




[ref](#)

Generative AI

Examples:

- Generative Adversarial Networks (GANs)




[ref](#)

Generative AI

Examples:

- Generative Adversarial Networks (GANs)




[ref](#)

Generative AI

Examples:

- Generative Adversarial Networks (GANs)




[ref](#)

Generative AI

Examples:

- Generative Adversarial Networks (GANs)




[ref](#)

Generative AI

Examples:

- Generative Adversarial Networks (GANs)
- **Variational Autoencoders (VAEs)**




[ref](#)

Foundational models

- We don't train independent models explicitly (say machine translation)

BERT


Encoder



GPT

Decoder

Machine translation




Foundational models

- We don't train independent models explicitly (say machine translation)
- Or document summarization

BERT

Encoder




GPT

Decoder

Document summarization

Foundational models

- We don't train independent models explicitly (say machine translation)
- Or document summarization
- **We train a Base model for multiple tasks jointly, and then fine tune for specific tasks**



Foundational models

- We don't train independent models explicitly (say machine translation)
- Or document summarization
- **We train a Base model for multiple tasks jointly, and then fine tune for specific tasks**




Figure 1.3.7

Training Neural Network Challenges

- Intractable gradients
 - Vanishing, and
 - Exploding gradients

Training Neural Network Challenges

- Intractable gradients
 - Vanishing, and
 - Exploding gradients
- Various normalizations
 - Input normalization (standard scalar)
 - Batch normalization
 - Layer normalization

Training Neural Network Challenges

- Intractable gradients
 - Vanishing, and
 - Exploding gradients
- Various normalizations
 - Input normalization (standard scalar)
 - Batch normalization
 - Layer normalization
- Controlling overfitting
 - Regularization
 - Early stopping
 - Drop out



QA